

MSA400: Technical Project 2

Calvin Smith, Kingsley Agyei Gyamfi

Chalmers University of Technology/University of Gothenburg

May 29, 2021

- Calvin: All the tasks (theory and coding) and everything concerning the Beta Binomial distribution.
- Kingsley: Logit normal distribution Task 1.

Contents

| | | |
|----------|---|-----------|
| 1 | Task 1 - Value at Risk | 3 |
| 1.1 | Beta Binomial | 3 |
| 1.2 | Logit-normal distribution | 4 |
| 2 | Task 2 - Choosing two sets of parameters | 4 |
| 2.1 | Beta-Binomial | 5 |
| 2.2 | Logit-normal distribution | 5 |
| 3 | Task 3 - Comparing models | 5 |
| 3.1 | Beta-Binomial | 5 |
| 3.2 | Logit-normal | 8 |
| 4 | Task 4 - Monte Carlo Simulation | 8 |
| 4.1 | Method (Beta-Binomial) | 8 |
| 4.2 | Results (Beta Binomial) | 9 |
| 5 | Task 5 - Estimation of parameters | 11 |
| 5.1 | Beta-Binomial | 11 |

1 Task 1 - Value at Risk

We are considering a static credit portfolio with 35 obligors where the size of each loan is 3 million SEK, the default losses are expected to be 60% and the individual default probability for each obligor is 4%. In this task, we have chosen two homogenous mixed binomial models, and we will derive formulas for obtaining the VaR using the Large Portfolio Approximation (LPA) approach.

The following theorem is central to our calculations:

Theorem 1 (*Large Portfolio Approximation for mixed binomial models*).
For large portfolios in a mixed binomial model, the distribution of the fractional number of defaults $\frac{N_m}{m}$ in the portfolio converges to the distribution of the random variable $p(Z)$ as $m \rightarrow \infty$, that is for any $x \in [0, 1]$ we have

$$\mathbb{P}\left[\frac{N_m}{m} \leq x\right] \rightarrow \mathbb{P}[p(Z) \leq x], \text{ when } m \rightarrow \infty. \quad (1)$$

The distribution $\mathbb{P}(p(Z) \leq x)$ is called the Large Portfolio Approximation (LPA) to the distribution of $\frac{N_m}{m}$.

From this, the following formula for LPA approximation of VaR will be used:

$$VaR_\alpha(L) \approx l \cdot m \cdot F^{-1}(\alpha) \quad (2)$$

where l is the default loss, m is the number of obligors and $F^{-1}(x)$ is the inverse of $F(x) = \mathbb{P}(p(Z) \leq x)$, where $p(Z)$ is the mixing distribution and Z is a random variable.

1.1 Beta Binomial

Letting $p(Z) = Z$ where Z is a beta distribution, $Z \sim \text{Beta}(a, b)$ we have that

$$F(x) = \mathbb{P}(Z \leq x) = \frac{1}{B(a, b)} \int_0^x z^{a-1} (1-z)^{b-1} dz. \quad (3)$$

where,

$$B(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}.$$

Here, $\Gamma(y)$ is the gamma function defined as $\Gamma(y) = \int_0^\infty t^{y-1} e^{-t} dt$.

$F(x)$ is by definition the CDF of a beta distribution, since for a continuous random variable X its Cumulative Distribution Function $F_X(x)$ can be expressed as

$$F_X(x) = \int_{-\infty}^x f_X(t) dt. \quad (4)$$

Thus, to obtain the value $F^{-1}(\alpha)$ for a given α we do not need to calculate an explicit formula, instead we can simply choose the x that satisfies $F(x) = \alpha$.

1.2 Logit-normal distribution

Let $p(Z)$ be a logit-normal distribution

$$F(x) = \frac{1}{1 + \exp(-u + \sigma Z)}, \quad (5)$$

where, $Z \sim N(0, 1)$, $p(Z) \in [0, 1]$. if $x \in [0, 1]$ then

$$p^{-1}(x) = \frac{1}{\sigma} \left(\ln\left(\frac{x}{1-x}\right) - u \right) \quad (6)$$

The mixing distribution $F(x) = \mathbb{P}[p(Z) \leq x] = \mathbb{P}[Z \leq P^{-1}(x)]$

for the logit-normal distribution is then giving by

$$F(x) = \mathbb{P}[Z \leq P^{-1}(x)] = \int_{-\infty}^{P^{-1}(x)} e^{-\frac{z^2}{2}} dz = N(P^{-1}(x)). \quad (7)$$

where $N(x)$ is the distribution function of the standard normal

Now to solve for VaR, $F^{-1}(\alpha)$ for a given α , we solve for x satisfying $F(x) = \alpha$

solving for X we have

$$X = \frac{\exp(\sigma F^{-1}(\alpha) + u)}{1 + \exp(\sigma F^{-1}(\alpha) + u)} \quad (8)$$

2 Task 2 - Choosing two sets of parameters

The second task consists of choosing two sets of parameter for each of the two models from Task 1. The first set should be chosen to provide a high default dependance and the second one low default dependance.

The default dependence of the obligors in a homogenous credit portfolio is measured by the correlation $Corr(X_i, X_j)$ between two pairs i, j in the portfolio. Since the portfolio is assumed to be homogenous the correlation between X_i and X_j is the same for all pairs as long as $i \neq j$. Instead of $Corr(X_i, X_j)$, we will write ρ_X to denote the correlation. Further, the correlation ρ_X in a mixed binomial model is given by the formula:

$$\rho_X = \frac{E[p(Z)^2] - \bar{p}^2}{\bar{p}(1 - \bar{p})} \quad (9)$$

where $\bar{p} = E[p(Z)] = \mathbb{P}[X_i = 1]$, is the default probability of each obligor.

2.1 Beta-Binomial

In the mixed binomial beta distribution we have that $p(Z) = Z$ where $Z \sim \text{Beta}(a, b)$. The expected value of a beta distributed random variable is

$$E[Z] = \frac{a}{a+b} \quad (10)$$

and the second moment is given by

$$E[Z^2] = \frac{(a+1)a}{(a+b+1)(a+b)}. \quad (11)$$

Since, the individual default probability for each obligor is 4% we have that

$$E[Z] = \frac{a}{a+b} = 0.04 = \bar{p}, \quad (12)$$

and thus we can rewrite the second moment as

$$E[Z^2] = \frac{(a+1)\bar{p}}{a+b+1}, \quad (13)$$

so inserting (8) and (9) in the formula for ρ_X we get

$$\rho_X = \frac{\frac{(a+1)\bar{p}}{a+b+1} - \bar{p}^2}{\bar{p}(1-\bar{p})} \quad (14)$$

which after some simplifications results in

$$\rho_X = \frac{1}{a+b+1}. \quad (15)$$

Thus, we can obtain a specified correlation ρ_X by choosing a and b that solves the system of equations:

$$\frac{a}{a+b} = 0.04 \quad (16)$$

$$\frac{1}{a+b+1} = \rho_X \quad (17)$$

From this, we can choose two sets of parameters $\{a, b\}$ for the Beta distribution, which results in two different models with the same mean $a/(a+b) = 0.04$, but with different correlations. The first set is $\{a = 0.01, b = 0.24\}$ which leads to a high correlation of $\rho_X = 0.8$. For the second set we chose $\{a = 0.36, b = 8.36\}$ which leads to a correlation $\rho_X = 0.1$.

2.2 Logit-normal distribution

3 Task 3 - Comparing models

3.1 Beta-Binomial

The value at risk estimates have been calculated using the formula (2) from Task 1. Further, the expected shortfall has been calculated using the following

| | a | b | μ | ρ_X |
|--------------|------|------|-------|----------|
| Beta Model 1 | 0.36 | 8.64 | 0.04 | 0.10 |
| Beta Model 2 | 0.01 | 0.24 | 0.04 | 0.80 |

Table 1: Two mixed beta binomial models with the same mean, but different correlations.

formula:

$$ES_\alpha(L) \approx \frac{l * m}{1 - \alpha} \int_\alpha^1 F^{-1}(u) du \quad (18)$$

where $F^{-1}(x)$ is the quantile function of our loss distribution.




Figure 1: Beta binomial probability density function and a binomial function as reference.




Figure 2: Beta binomial probability function and a binomial function as reference. Zoomed in to capture the difference in the tails.

Figure 1 and 2 shows the probability density functions of our two mixed binomial beta models and a standard binomial probability density function as a reference. As we can see, the most noticeable difference is how the model with a high dependance between obligors acts around the tails, meaning that the number of defaults gets pushed to the extremes. Since the probability of an individual default is fairly small ($p = 0.04$) but dependance between obligors in this model is very high, we have a scenario where there is a high probability of all obligors defaulting and all obligors not defaulting, while the probability of some defaulting and some not is extremely low. The low dependancy model is on the other hand more centered around the mean. For example, the probability of only 1 default in the high dependency model is 0.009 and in the low dependency model it is 0.16. Conversely, the probability of 35 defaults in the high dependency model is 0.015, while in the low dependency model it is 0.000000000285. This also corresponds well to the results from table 2 and 3 where we see a much higher value at risk and expected shortfall for the high dependency model compared to the low dependency model.

| α | 95% | 99% | 99.9 |
|----------|-------|-------|-------|
| VaR | 10.70 | 18.31 | 27.91 |
| ES | 15.36 | 22.53 | 31.35 |

Table 2: VaR and ES estimates for Beta Model 1.

| α | 95% | 99% | 99.9 |
|----------|-------|-------|-------|
| VaR | 13.27 | 62.81 | 63.00 |
| ES | 51.26 | 62.96 | 63.00 |

Table 3: VaR and ES estimates for Beta Model 2.

3.2 Logit-normal

4 Task 4 - Monte Carlo Simulation

4.1 Method (Beta-Binomial)

As a comparison to the LPA approximations we want to use Monte Carlo simulations to generate an empirical loss distribution in order to calculate VaR and ES. The outline of the Monte Carlo algorithm works as follows:

1. Let n be the number of simulations.
2. for $j = 1, 2, \dots, n$:
 - Generate a sample z_j from the random variable $Z \sim \text{Beta}(a, b)$.
 - Generate m samples u_1, u_2, \dots, u_m from a uniformly distributed random variable $U \sim \text{Uniform}(0, 1)$, where m is the number of obligors.
3. for $i = 1, 2, \dots, m$:

We create a vector of defaults X_i by

$$X_i = \begin{cases} 1 & \text{if } u_i \leq z_j \\ 0 & \text{otherwise} \end{cases} \quad (19)$$

4. Compute $L_m^{(j)} = l \sum_{i=1}^m X_i$, where $L_m^{(j)}$ is the total loss from simulation j and l is the default loss, constant for all obligors.

From the simulated empirical loss distribution $\{L_m^{(j)}\}_{j=1}^n$ we can calculate the empirical VaR by finding the α -quantile of our distribution, which is

$$\widehat{VaR}_\alpha(L) = q_\alpha(\{L_m^{(j)}\}_{j=1}^n) = \{L_m^{(j)}\}_{j=n*\alpha} \quad (20)$$

and

$$\widehat{ES}_\alpha(L) = \frac{\sum_{k=n*\alpha}^n \{L_m^{(j)}\}_k}{n(1 - \alpha)}. \quad (21)$$

In the implementation of the Monte Carlo simulation we have used $n = 1000000$.

4.2 Results (Beta Binomial)

The results of the VaR and ES from the Monte Carlo simulations are presented in tables 4 and 5. Comparing with the LPA approximations from table 2 and 3 we can see that the results are quite similar for both models, concluding that in this case LPA approximation works fairly well.

| α | 95% | 99% | 99.9% |
|----------|-------|-------|-------|
| VaR | 12.60 | 19.80 | 30.60 |
| ES | 16.98 | 23.76 | 33.61 |

Table 4: Results from the MC simulation with Beta Model 1.

| α | 95% | 99% | 99.9% |
|----------|-------|-------|-------|
| VaR | 12.60 | 63.00 | 63.00 |
| ES | 46.92 | 63.00 | 63.00 |

Table 5: Results from MC simulation with Beta Model 2.




Figure 3: Histogram of the MC loss distribution - Beta Model 1.

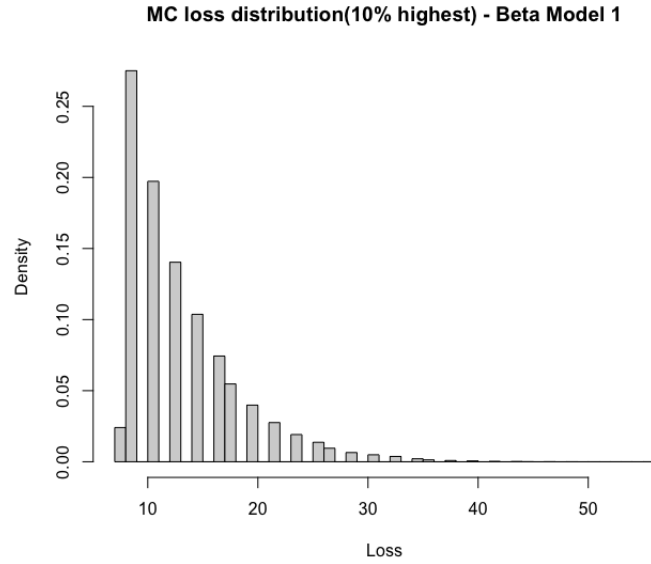


Figure 4: Histogram of the MC loss distribution(highest 10%) - Beta Model 1

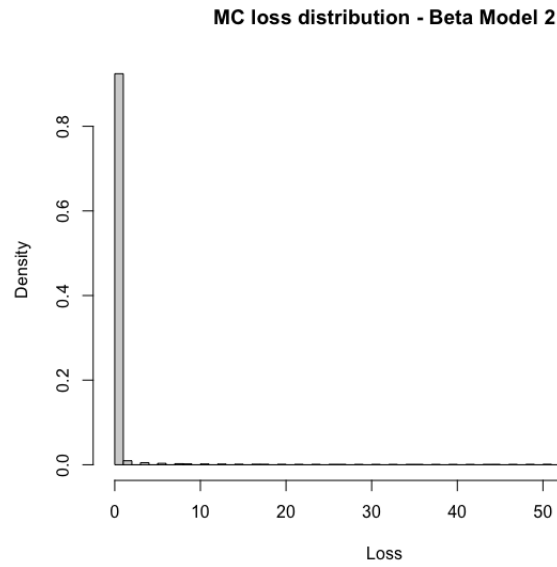


Figure 5: Histogram of the MC loss distribution - Beta Model 2.




Figure 6: Histogram of the MC loss distribution(highest 10%) - Beta Model 2

5 Task 5 - Estimation of parameters

5.1 Beta-Binomial

In the case of the Beta-Binomial framework, the easiest (and most straight-forward) way of estimating the parameters of our model is probably using the method of moments, since the Bayesian framework of using a Binomial likelihood and a Beta conjugate prior is well defined and gives us a nice formula to compute the first two moments of our Beta-Binomial distribution. The first two moments of the Beta-Binomial distribution are given by

$$\mu_1 = \frac{ma}{a+b} \quad (22)$$

$$\mu_2 = \frac{ma(m(1+a)+b)}{(a+b)(1+a+b)} \quad (23)$$

where m is the number of obligors and a, b are the parameters to be estimated. So, assuming that we have some sample data of defaults $\mathbf{X} = (X_1, X_2, \dots, X_n)$ where each X_i represents the number of defaults (for a given number of obligors m) over a set time period T , we can use the sample moments

$$\hat{\mu}_1 = \frac{1}{n} \sum_{i=1}^n X_i = k_1$$

$$\hat{\mu}_2 = \frac{1}{n} \sum_{i=1}^n X_i^2 = k_2$$

and setting them equal to (22) and (23), we solve for a and b and get the estimates:

$$\hat{a} = \frac{nk_1 - k_2}{n(k_2/k_1 - k_1 - 1) + k_1} \quad (24)$$

$$\hat{b} = \frac{(n - k_1)(n - k_2/k_1)}{n(k_2/k_1 - k_1 - 1) + k_1}. \quad (25)$$

Another approach is of course to use Maximum Likelihood to estimate the parameters of our model. However, finding closed form expressions for the MLE-estimates in a Beta-Binomial model is not obvious (at least not for me) and one should use a built in package in some programming language (for example the VGAM package in R) that can solve the problem numerically. From reading about the **betabinomial**-function in VGAM, it uses the Fisher Scoring algorithm to solve the Maximum Likelihood equations.

Appendix

R code - Beta Binomial implementation

```
#####
##### TASK 1 #####
#####

OBLIGORS <- 35
LOAN <- 3
DEFAULT_LOSSES <- 0.6
DEFAULT_PROBABILITY <- 0.04

a <- 0.2
b <- 4.59

VaR_beta <- LOAN*DEFAULT_LOSSES*OBLIGORS*qbeta(alpha,a,b)

#####
##### TASK 2 and 3 #####
#####
library("extraDistr")

DEFAULTS <- 0:35

# low correlation (corr = 0.1)
a1 <- 0.36
b1 <- 8.64

mod1 <- c(a1,b1,a1/(a1+b1),0.1)

# high correlation (corr = 0.8)
a2 <- 0.01
b2 <- 0.24

mod2 <- c(a2,b2,a2/(a2+b2),0.8)

summary <- rbind(mod1,mod2)
rownames(summary) <- c("Model 1","Model 2")
colnames(summary) <- c("a","b","p","rho")

# Fitting density functions
```

```

betamodel_1 <- dbbinom(DEFAULTS,OBLIGORS,alpha = a1,beta = b1)

betamodel_2 <- dbbinom(DEFAULTS,OBLIGORS,alpha = a2,beta = b2)

binomial <- dbinom(DEFAULTS,OBLIGORS,0.04)

plot(DEFAULTS,betamodel_1,"o",col = "green",bg = "green",lwd = 1,
ylim = c(0,0.2),main = "Beta binomial" ,xlab = "Number of defaults",ylab = "Probability")
lines(DEFAULTS,betamodel_2,"o",col = "red",bg = "red",lwd = 1)
lines(DEFAULTS,binomial,"o",col = "blue",bg="blue",lwd = 1)

color.names <- c("green","red","blue")
legend("top",c("Mixed binomial 35, Beta(0.36,8.64)",
"Mixed binomial 35, Beta(0.01,0.24)","Binomial (35,0.04)"),fill = color.names)

# Value at Risk
quantiles <- c(0.95,0.99,0.999)

VaR_beta1 <- LOAN*DEFAULT_LOSSES*OBLIGORS*qbeta(quantiles,a1,b1)

VaR_beta2 <- LOAN*DEFAULT_LOSSES*OBLIGORS*qbeta(quantiles,a2,b2)

# Expected Shortfall
ES_model1 <- numeric(3)
ES_model2<- numeric(3)

for (i in 1:length(quantiles)){
  ES_model1[i] <- ((LOAN*OBLIGORS*DEFAULT_LOSSES)/(1-quantiles[i]))*integrate(qbeta,
lower = quantiles[i],upper = 1,shape1 = a1,shape2=b1)$value
  ES_model2[i] <- ((LOAN*OBLIGORS*DEFAULT_LOSSES)/(1-quantiles[i]))*integrate(qbeta,
lower = quantiles[i],upper = 1,shape1 = a2,shape2=b2)$value
}

# Making tables for LaTeX
model1_var_es <- rbind(VaR_beta1,ES_model1)
rownames(model1_var_es) <- c("VaR","ES")
colnames(model1_var_es) <- c("95%","99%","99.9")

model2_var_es <- rbind(VaR_beta2,ES_model2)
rownames(model2_var_es) <- c("VaR","ES")
colnames(model2_var_es) <- c("95%","99%","99.9")

library(xtable)

xtable(model1_var_es)

```

```

xtable(model2_var_es)

#####
##### TASK 4: MONTE CARLO SIMULATION #####
#####

# function for Monte Carlo simulation
monte_carlo <- function(nsim,a,b,obligors,loan,default_losses){
  L <- numeric(nsim)
  for (j in 1:nsim){
    Z <- rbeta(1,shape1 = a,shape2 = b)
    U <- runif(obligors,0,1)
    X <- numeric(obligors)
    for (i in 1:obligors){

      if (U[i] <= Z){
        X[i] <- 1
      }
      else{
        X[i] <- 0
      }
    }
    L[j] <- sum(X)*loan*default_losses
  }
  return(L)
}

n <- 1000000

# Beta Binomial model 1, VaR
L_model1 <- monte_carlo(n,a1,b1,OBLIGORS,LOAN,DEFAULT_LOSSES)
MC_VaR_model1 <- quantile(L_model1,quantiles)

# Beta Binomial model 2, VaR
L_model2 <- monte_carlo(n,a2,b2,OBLIGORS,LOAN,DEFAULT_LOSSES)
MC_VaR_model2 <- quantile(L_model2,quantiles)

# function for calculating Expected Shortfall

ES <- function(vaR,losses){

  expectedShortfall <- numeric(length(vaR))

  for (i in 1:length(vaR)){

```

```

        expectedShortfall[i] <- sum(losses[which(losses >= vaR[i])])/
        (length(which(losses >= vaR[i])))
    }
    return(expectedShortfall)
}

# Beta Binomial Model 1, ES
MC_ES_model1 <- ES(MC_VaR_model1,L_model1)

# Beta Binomial Model 1, ES
MC_ES_model2 <- ES(MC_VaR_model2,L_model2)

# Creating tables for LaTeX
summaryMC_1 <- rbind(MC_VaR_model1,MC_ES_model1)
rownames(summaryMC_1) <- c("VaR","ES")

summaryMC_2 <- rbind(MC_VaR_model2,MC_ES_model2)
rownames(summaryMC_2) <- c("VaR","ES")

# plots of MC distributions

hist(L_model1,freq = FALSE, breaks = 50,
main = "MC loss distribution - Beta Model 1", xlab = "Loss")

hist(L_model2,freq = FALSE, breaks = 50,
main = "MC loss distribution - Beta Model 2", xlab = "Loss")

hist(sort(L_model1)[900000:1000000],freq = FALSE,breaks = 50,
main ="MC loss distribution(10% highest) - Beta Model 1", xlab = "Loss")

hist(sort(L_model2)[900000:1000000],freq = FALSE, breaks = 50,
main = "MC loss distribution(10% highest - Beta Model 2", xlab = "Loss")

```